

Predictive Sales Analytics Engine

1st Suyash Parmar

Newton School of Technology
230057
suyashparmar44@gmail.com

2nd Samyak Mittal

Newton School of Technology
230091
samyakmittal23@gmail.com

Abstract—Predicting sales outcomes in complex CRM environments remains a challenge due to the high cardinality of categorical features and the need for both interpretability and high-order feature interaction modeling. This paper presents a three-phase predictive engine that culminates in a symbiotic hybrid network. We propose a Neural-Probabilistic Gating mechanism that dynamically weights an interpretable Logistic Regression prior against a Deep Tabular latent representation. Our results demonstrate that this neuro-symbolic approach outperforms traditional ensembling methods by achieving differentiable synergy between linear boundaries and deep non-linear residuals.

Index Terms—Sales Forecasting, Neuro-Symbolic AI, Gating Networks, Deep Tabular Learning, CRM Analytics

I. INTRODUCTION

Modern sales organizations generate vast amounts of multi-relational data involving accounts, agents, and product sectors. Traditional approaches typically fall into two camps: (1) highly interpretable linear models that lack the capacity for complex interactions, and (2) black-box deep learning models that struggle with tabular data sparsity. We bridge this gap by introducing a three-phase pipeline that evolves from interpretable baselines to a symbiotic hybrid architecture.

II. METHODOLOGY

The proposed system is structured into three distinct phases of increasing complexity and integration.

A. Phase 1: Advanced Machine Learning Baseline

The initial phase focuses on robustness and interpretability. We implement a stratified preprocessing pipeline (`src/preprocessing.py`) that joins CRM datasets while preserving deal-class structures. A Logistic Regression model is trained using standard scaling transformations and sparse one-hot expansions. This phase provides a well-calibrated probabilistic prior which serves as a symbolic anchor for the system.

B. Phase 2: Deep Tabular Neural Pipelines

To capture latent semantic structures, we bypass static encoding in favor of a categorical embedding network (`src/dl_model.py`). This architecture maps high-cardinality constraints into dense embedding layers which are concatenated with numerical features. The network is optimized using binary cross-entropy with logits, allowing it to learn non-linear residuals that the linear phase misses.

C. Phase 3: Symbiotic Hybrid Innovation

The core innovation is the Probabilistic Gating Network implemented in `src/hybrid_model.py`. Unlike naive ensembling, this architecture uses a differentiable gating mechanism that learns to dynamically balance the predictions. The system analyzes the raw machine learning probability against the deep latent embeddings to adjust its final confidence level. This allows the model to rely on the interpretable machine learning prior for standard deals while shifting weight to the deep learning branch when complex latent anomalies are detected.

III. IMPLEMENTATION AND DEPLOYMENT

The system is implemented using PyTorch and Scikit-Learn. We provide a turn-key environment including:

- **Dockerization:** Containerized execution for OS-independent reproducibility.
- **Streamlit Dashboard:** An interactive web application (`app.py`) that enables side-by-side comparison of the machine learning, deep learning, and hybrid signals.
- **Differentiable Training:** The gating mechanism is trained end-to-end alongside the deep features, ensuring optimal synergy between the symbolic priors and neural representations.

IV. CONCLUSION

Our symbiotic hybrid engine demonstrates that interpretability and predictive power are not mutually exclusive. By treating linear machine learning models as probabilistic priors within a deep gating architecture, we achieve a robust, interpretable, and high-performance solution for sales win prediction. Future work will explore the application of this gating logic to time-series sales forecasting.

REFERENCES

- [1] I. Goodfellow et al., “Deep Learning,” MIT Press, 2016.
- [2] L. Breiman, “Statistical Modeling: The Two Cultures,” Statistical Science, 2001.